

Load the Dataset

```
import pandas as pd
import numpy as np
from matplotlib import rcParams
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv('/content/winequality-red.csv')
df.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

Data preprocessing including visualization

```
df.isnull().any()

fixed acidity      False
volatile acidity   False
citric acid        False
residual sugar     False
chlorides          False
free sulfur dioxide False
total sulfur dioxide False
density            False
pH                 False
sulphates          False
alcohol            False
quality            False
dtype: bool
```

```
df.isnull().sum()

fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
sulphates          0
alcohol            0
quality            0
dtype: int64
```

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   fixed acidity       1599 non-null   float64
 1   volatile acidity    1599 non-null   float64
 2   citric acid         1599 non-null   float64
 3   residual sugar      1599 non-null   float64
 4   chlorides           1599 non-null   float64
 5   free sulfur dioxide 1599 non-null   float64
 6   total sulfur dioxide 1599 non-null   float64
 7   density             1599 non-null   float64
 8   pH                  1599 non-null   float64
 9   sulphates           1599 non-null   float64
10   alcohol             1599 non-null   float64
11   quality             1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```


```
df.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	su
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996747	3.311113	0.016101
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001887	0.154386	0.000254
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	0.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995600	3.210000	0.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996750	3.310000	0.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997835	3.400000	0.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.003690	4.010000	0.000000

```
df.shape
```

(1599, 12)

```
sns.boxplot(df['fixed acidity'])
```



```
sns.distplot(df['fixed acidity'])
```


```
<ipython-input-12-52a4a49dcd39>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
sns.boxplot(df['volatile acidity'])
```

```
<Axes: >
```



```
sns.distplot(df['volatile acidity'])
```

```
<ipython-input-14-6077730c287e>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['volatile acidity'])
```

```
<Axes: xlabel='volatile acidity', ylabel='Density'>
```



```
sns.boxplot(df['citric acid'])
```

<Axes: >



```
sns.distplot(df['citric acid'])
```

```
<ipython-input-16-1324198882c2>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['citric acid'])
```

```
<Axes: xlabel='citric acid', ylabel='Density'>
```



```
sns.boxplot(df['residual sugar'])
```

<Axes: >



```
sns.distplot(df['residual sugar'])
```

```
<ipython-input-18-17c4014efccf>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see


<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['residual sugar'])  
<Axes: xlabel='residual sugar', ylabel='Density'>
```



```
sns.boxplot(df['chlorides'])
```

```
<Axes: >
```



```
sns.distplot(df['chlorides'])
```

```
<ipython-input-20-fdc4bb1ed131>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['chlorides'])
<Axes: xlabel='chlorides', ylabel='Density'>
```



```
sns.boxplot(df['free sulfur dioxide'])
```

```
<Axes: >
```



```
sns.distplot(df['free sulfur dioxide'])
```


```
<ipython-input-24-3dee0624d434>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['free sulfur dioxide'])
<Axes: xlabel='free sulfur dioxide', ylabel='Density'>
```



```
sns.boxplot(df['total sulfur dioxide'])
```



```
sns.distplot(df['total sulfur dioxide'])
```

<ipython-input-26-a53ba4eac084>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['total sulfur dioxide'])
<Axes: xlabel='total sulfur dioxide', ylabel='Density'>
```



```
sns.boxplot(df['density'])
```

<Axes: >

```
sns.distplot(df['density'])
```

<ipython-input-28-cffea316cede>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['density'])
```

<Axes: xlabel='density', ylabel='Density'>



```
sns.boxplot(df['pH'])
```

<Axes: >



```
sns.distplot(df['pH'])
```



```
<ipython-input-30-d020e64af2d2>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see


<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['pH'])  
<Axes: xlabel='pH', ylabel='Density'>
```




```
sns.boxplot(df['sulphates'])
```

```
<Axes: >
```



```
sns.distplot(df['sulphates'])
```

```
<ipython-input-32-3a090c5692ad>:1: UserWarning:
sns.boxplot(df['alcohol'])
```



```
sns.distplot(df['alcohol'])
```


```
<ipython-input-34-570de8ff0310>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with
similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751
```

```
sns.distplot(df['alcohol'])
<Axes: xlabel='alcohol', ylabel='Density'>
```



```
sns.boxplot(df['quality'])
```

<Axes: >



```
sns.distplot(df['quality'])
```

```
<ipython-input-36-e9b2f3ff6ab5>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

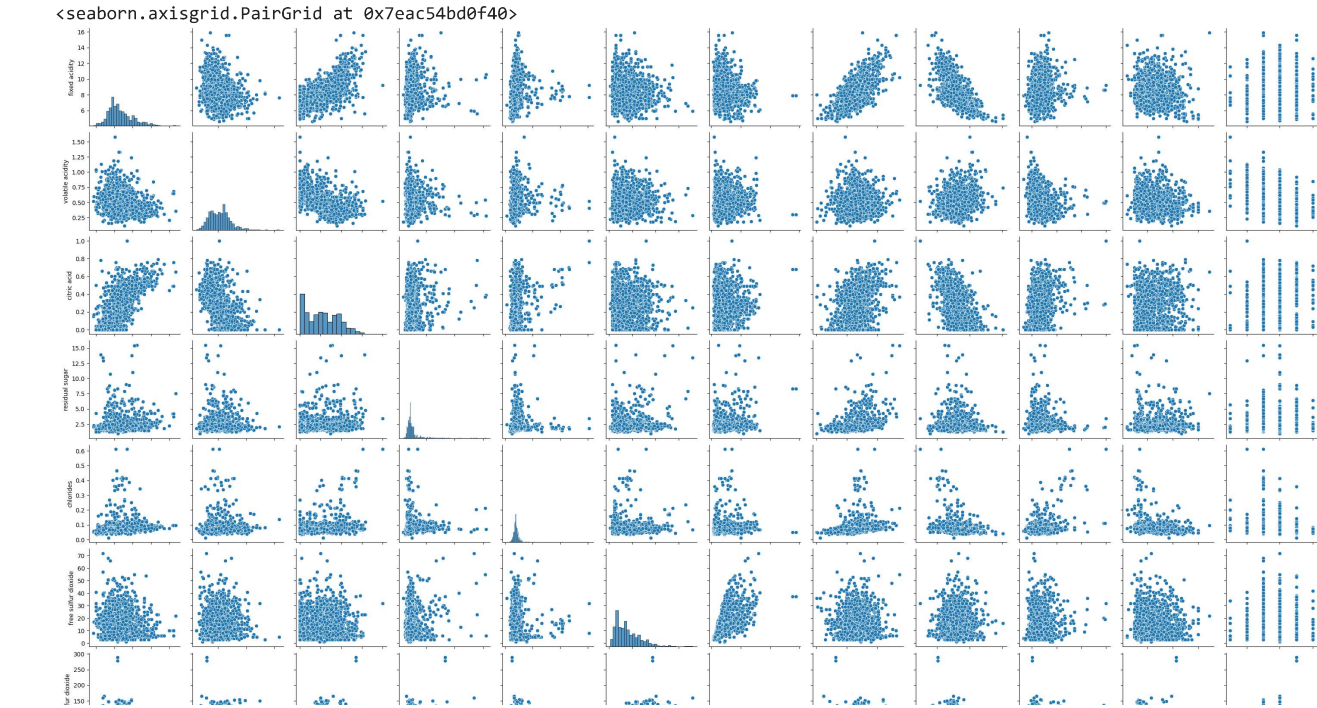
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['quality'])
```

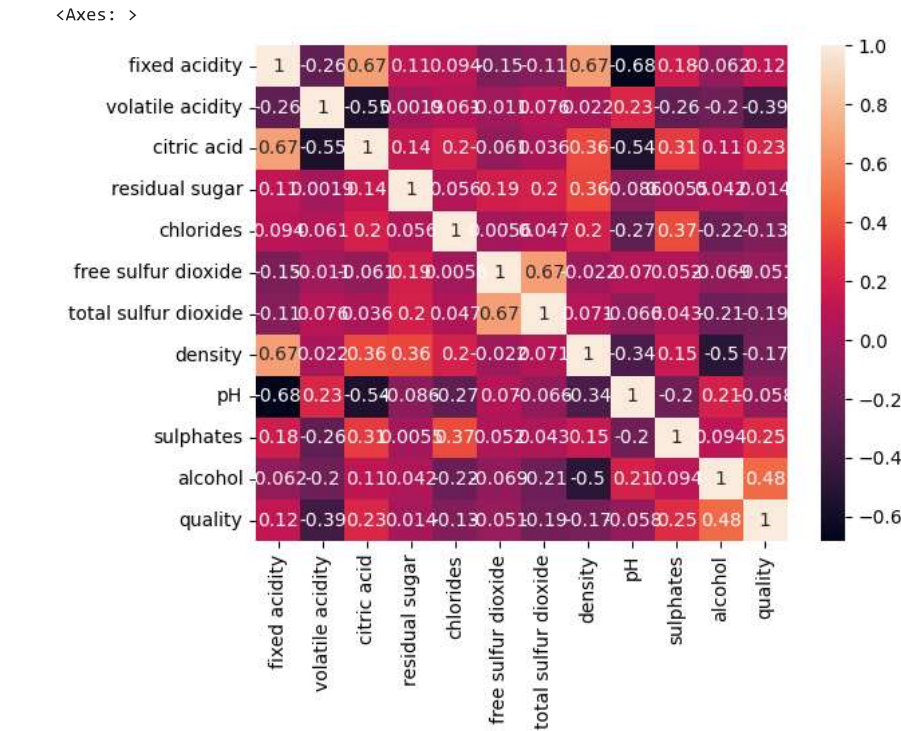
```
<Axes: xlabel='quality', ylabel='Density'>
```



```
sns.pairplot(df)
```



```
sns.heatmap(df.corr(),annot = True)
```



```
X = df.drop(columns = ['quality'], axis = 1)
X.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4

y = df.quality
y.head()

0	5
1	5
2	5

```
3 6
4 5
Name: quality, dtype: int64
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train, y_test = train_test_split(X,y, test_size= 0.2, random_state = 0)
```

```
X_train.shape
(1279, 11)
```

```
X_test.shape
(320, 11)
```

```
y_train.shape
(1279,)
```

```
y_test.shape
(320,)
```

Machine Learning Model building

```
from sklearn.linear_model import LinearRegression, LogisticRegression
lr = LinearRegression()
lor = LogisticRegression()
```

```
lr.fit(X_train,y_train)
```

```
LinearRegression()
```

```
lor.fit(X_train,y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: Conver
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
LogisticRegression
```

```
LogisticRegression()
```

```
y_predict_lr = np.round(lr.predict(X_test))
y_predict_lr
```

```
array([6., 5., 7., 5., 6., 5., 5., 6., 5., 5., 5., 5., 6., 5., 6., 6., 7.,
        6., 6., 5., 6., 5., 6., 6., 5., 5., 5., 6., 5., 6., 6., 6., 5.,
        6., 5., 5., 6., 6., 5., 6., 7., 7., 6., 5., 5., 6., 5., 6., 5.,
        5., 6., 6., 6., 5., 5., 5., 7., 5., 5., 6., 6., 6., 5., 6., 5., 6.,
        6., 6., 5., 5., 5., 6., 6., 6., 5., 5., 6., 6., 6., 6., 6.,
        5., 5., 5., 5., 5., 6., 5., 6., 5., 6., 5., 5., 6., 7., 6.,
        6., 6., 5., 6., 5., 6., 5., 6., 5., 6., 5., 6., 6., 6., 6.,
        6., 5., 6., 5., 5., 6., 6., 5., 5., 6., 6., 5., 5., 6., 6., 5.,
        6., 5., 6., 5., 5., 5., 5., 6., 6., 6., 6., 6., 5., 6., 6., 5.,
        6., 6., 5., 7., 6., 6., 6., 7., 6., 5., 5., 7., 5., 6., 7., 5., 6.,
        6., 5., 6., 6., 6., 5., 5., 5., 5., 5., 5., 5., 5., 6., 5., 5.,
        5., 5., 5., 6., 6., 5., 6., 6., 5., 5., 5., 5., 6., 6., 6., 5.,
        6., 6., 6., 6., 6., 5., 5., 5., 5., 6., 5., 6., 5., 5., 5.,
        6., 6., 6., 7.,
        7., 6., 5., 5., 5., 5., 6., 5., 6., 5., 5., 6., 5., 5., 5., 6.,
        6., 5., 5., 5., 6., 5., 7., 5., 6., 5., 5., 5., 5., 6., 6., 6.,
        6., 6., 6., 6., 6., 5., 7., 6., 5., 7., 6., 6., 6., 5., 6., 5., 6.,
        6., 6., 5., 6., 5., 5., 6., 6., 5., 5., 5., 6., 5., 5., 6., 6.,
        5., 5., 6., 5., 6., 6., 5., 5., 5., 7., 6., 6., 5., 6.]
```

```
y_predict_lor = np.round(lor.predict(X_test))
y_predict_lor
```

```
array([6, 5, 6, 5, 6, 5, 5, 6, 5, 5, 5, 5, 6, 5, 6, 6, 7, 6, 6, 5, 6, 5,
        6, 6, 5, 5, 5, 6, 5, 6, 6, 6, 6, 5, 6, 6, 5, 5, 6, 6, 5, 6, 6, 7,
        6, 5, 6, 6, 5, 6, 5, 5, 6, 6, 5, 5, 5, 5, 6, 5, 5, 6, 6, 6, 5, 6,
```

```
6, 6, 6, 6, 5, 6, 5, 6, 6, 6, 5, 5, 5, 6, 6, 5, 6, 6, 6, 5, 6, 5,
5, 5, 5, 5, 6, 5, 6, 5, 6, 5, 5, 6, 7, 6, 6, 6, 6, 5, 6, 5, 6, 5,
6, 5, 6, 5, 6, 5, 6, 7, 6, 6, 5, 6, 6, 5, 6, 6, 5, 5, 6, 6, 5, 5,
6, 6, 6, 5, 6, 5, 6, 5, 6, 5, 6, 5, 5, 5, 6, 6, 6, 5, 6, 6, 5, 6,
5, 6, 5, 5, 6, 6, 6, 5, 6, 5, 6, 5, 6, 6, 5, 6, 6, 5, 5, 6, 6, 6,
7, 7, 6, 5, 5, 6, 5, 6, 7, 5, 6, 6, 5, 6, 6, 6, 5, 5, 5, 5, 5, 5,
5, 5, 5, 6, 5, 5, 6, 5, 5, 6, 6, 5, 6, 6, 5, 7, 5, 5, 6, 6, 5, 5,
5, 6, 6, 6, 5, 6, 6, 6, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 5, 5,
5, 5, 6, 5, 5, 5, 5, 6, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 6, 5, 5,
5, 5, 5, 5, 5, 6, 6, 5, 6, 6, 6, 6, 6, 5, 6, 6, 5, 6, 6, 5,
6, 5, 6, 6, 6, 6, 6, 5, 5, 6, 5, 5, 5, 5, 6, 6, 6, 5, 5,
5, 5, 6, 6, 5, 5, 5, 7, 6, 6, 5, 6])
```

Evaluate the model

```
from sklearn import metrics
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
as1 = accuracy_score(y_test, y_predict_lr)
as1

0.63125
```

```
as2 = accuracy_score(y_test, y_predict_lor)
as2

0.625
```



```
r2s1 = metrics.r2_score(y_test,y_predict_lr)
r2s1

0.20846127601501196
```

```
r2s2 = metrics.r2_score(y_test, y_predict_lor)
r2s2

0.1647901740020471
```

```
pd.crosstab(y_test,y_predict_lr)
```

col_0	5.0	6.0	7.0	
quality				
3	2	0	0	
4	6	5	0	
5	97	38	0	
6	37	98	7	
7	1	19	7	
8	0	1	2	

```
pd.crosstab(y_test, y_predict_lor)
```

col_0	5	6	7	
quality				
3	2	0	0	
4	7	4	0	
5	101	33	1	
6	42	96	4	
7	2	22	3	
8	0	2	1	

```
classification_report(y_test, y_predict_lr)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-sc
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-sc
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-sc
_warn_prf(average, modifier, msg_start, len(result))

classification_report(y_test,y_predict_lor)

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-sc
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-sc
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-sc
_warn_prf(average, modifier, msg_start, len(result))
'
precision    recall  f1-score   support\n
 4      0.00      0.00      0.00      11\n
 5      0.66      0.75      0.70      135\n
 6      0.68      0.64      0.66      142\n
 7      0.33      0.11      0.17      27\n
 8      0.00      0.00      0.00      320\n
00      3\n\n
accuracy      0.58      0.62      0.59      320\n'
macro avg      0.27      0.26      0.25      3
20\nweighted avg
```

Test with random observation

```

random_input_1 = np.round(lr.predict([[9.0,0.60,0.05,1.8,0.087,11.0,25.0,0.8888,8.90,0.90,7.4]]))
random_input_1

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression
warnings.warn(
array([7.])

random_input_2 = np.round(lr.predict([[9.4,0.90,0.10,3.0,0.155,12.0,90.0,0.8888,2.50,0.68,9.6]]))
random_input_2

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression
warnings.warn(
array([8.])
```